

# Traditional Machine Learning

Theory and Implementation



UNIVERSITY OF  
CALGARY

# Outline

---

- Learning Goals
- Traditional ML Concepts
- Popular ML algorithms
- Model Evaluation
- ML best practices
- Summary



# Learning Goals

---

- **Understand** how machines learn patterns from data
- **Learn** the foundations of key ML algorithms
- **Apply** traditional ML methods to real-world problems
- **Evaluate** model performance using appropriate metrics
- **Recognize** when to use traditional ML vs. deep learning

# What is Machine Learning?

---

*"Machine Learning is the field of study that gives computers the ability to learn without being explicitly programmed."*

— Arthur Samuel (1959)

- **Modern Perspective**

- ML algorithms learn patterns from data to make predictions or decisions
- Focus on building models that improve performance through experience
- Bridges statistics, computer science, and domain expertise
- Examples: Email spam detection, Medical diagnosis assistance

# Machine Learning Paradigms

## Artificial Intelligence

### Machine Learning

#### Supervised Learning

- Learns from labeled data
- Has target/output variable
- Types:
  - Regression(continuous)
  - Classification (discrete)

#### Unsupervised Learning


- Learns from unlabeled data
- Finds hidden patterns
- Types:
  - Clustering
  - Dimensionality reduction

#### Reinforcement Learning

- Learn through feedback
- No predefined data learn from environment
- Examples:
  - Game AI
  - Robotics

## Deep Learning

# Machine Learning Paradigms



# Traditional ML vs Deep Learning

---


## Traditional ML

- Requires feature engineering
- Works well with small datasets
- More interpretable
- Faster training
- Less computational resources
- Examples: Linear Regression, SVM, Random Forest

## Deep Learning

- Automatic feature learning
- Requires large datasets
- Less interpretable
- Longer training time
- GPU-intensive
- Examples: CNN, RNN, Transformers

# Machine Learning Workflow



# Feature Engineering

*The art of transforming raw data into meaningful features*

## Data Preprocessing

- Handle missing values
  - Removal, mean/median/mode imputation
- Handle outliers
  - Removal, capping, transformation
- Feature scaling
  - Normalization, standardization
- Splitting data
  - Train/Test/Validation

## Feature Transformation

- Encode categorical
  - Label encoding, one-hot encoding
- Polynomial features
- Log, square root, or power transform
- Interaction features
  - Combining two or more features
- Dimensionality reduction
  - PCA, LDA

## Feature Selection

- Filter methods
- Recursive Feature Elimination (RFE)
- Removing low-variance features
- Regularization methods
  - Lasso, Ridge
- Feature importance from models



Good features often matter more than the algorithm!

# Train/Test/Validation Split

---

- **Why Split Data?**
  - To evaluate how well our model generalizes to unseen data
- **Common Split Strategy:**
  - **Training Set (60-70%):** Learn patterns from data
  - **Validation Set (15-20%):** Tune hyperparameters
  - **Test Set (15-20%):** Final unbiased evaluation

 **Critical Rule: NEVER train on test data!**

# Linear Regression

- Goal: Models linear relationship between features and target

## Model Formulation


$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n$$

Objective: Minimize Sum of Squared Errors

$$\text{Loss}(\beta) = \sum (y_i - \hat{y}_i)^2$$

✓ Pros: Simple, low cost, fast, interpretable

✗ Cons: Assumes linearity, Sensitive to outliers



Applications: Price prediction • Sales forecasting • Trend analysis • Risk assessment

# Logistic Regression


- Goal: Classify data into discrete categories (typically binary: 0 or 1)

**Sigmoid Function:**  
 $\sigma(z) = 1 / (1 + e^{-z})$

where  $z = \beta_0 + \beta_1 x_1 + \dots + \beta_n x_n$   
Output:  $P(y=1|x) \in [0, 1]$

**Loss Function:**  
Cross-Entropy Loss

$$L = -[y \log(\hat{y}) + (1-y) \log(1-\hat{y})]$$



✓ Pros: Probabilistic, interpretable

✗ Cons: Linear boundary

Applications: Email Spam Detection • Medical diagnosis • Credit risk

# Decision Trees


- Goal: Create a tree-like model of decisions for classification or regression
  - Splits data recursively based on feature values
    - Internal nodes = feature tests
    - Branches = test outcomes
    - Leaves = predicted class or value

- **Splitting Criteria**

- **Gini Impurity:** Measures node impurity
  - $\text{Gini} = 1 - \sum(p_i)^2$
- **Entropy (Information Gain):** Measures information
  - $\text{Entropy} = -\sum p_i \log_2(p_i)$

✓ Pros: Interpretable, Handles non-linear data

✗ Cons: Prone to overfitting, Unstable to small data changes



# Random Forests

---

- Goal: Ensemble of decision trees - majority vote wins
- **How it Works?**
  - **Bootstrap Sampling:** Random samples with replacement
  - **Random Feature Selection:** Each split uses subset of features
  - **Train Multiple Trees:** Each tree on different sample
  - **Aggregate Predictions:** Majority vote or average

✓ Pros: More accurate, less overfitting

✗ Cons: Less interpretable

# K-Means Clustering


- Goal: Partition data into K clusters

- **Steps:**

1. **Initialize:** Choose K random centroids
2. **Assignment:** Assign each point to nearest centroid
3. **Update:** Recalculate centroids as mean of assigned points
4. **Repeat:** Until convergence

✓ Pros: Fast, scalable

✗ Cons: Must choose K



# When to Use Which Algorithm?

Algorithm	Best For	Pros	Cons
<b>Linear Regression</b>	Linear relationships	Fast, interpretable	Assumes linearity
<b>Logistic Regression</b>	Binary classification	Simple, probabilistic	Linear boundary
<b>Decision Trees</b>	Non-linear, interpretability	Easy to visualize	Overfitting
<b>Random Forest</b>	Complex data	Robust, accurate	Less interpretable
<b>SVM</b>	High-dimensional	Effective in high dims	Slow on large data
<b>K-Nearest Neighbors</b>	Instance-based learning	Simple, no training	Slow prediction, memory intensive
<b>K-Means</b>	Segmentation	Simple, fast	Needs k specified

# Model Evaluation

---

## *Measuring and Validating Model Performance*

### Why Evaluate?

- Assess model performance
- Compare different algorithms
- Detect overfitting/underfitting
- Make informed decisions

### What to Evaluate?

- Prediction accuracy
- Generalization ability
- Computational efficiency
- Model interpretability

### How to Evaluate?

- Appropriate metrics
- Cross-validation
- Train/test splits
- Confusion matrix analysis

# Confusion Matrix

Confusion Matrix

		Predicted	
		Positive	Negative
Actual	Positive	TP	FP
	Negative	FN	TN

## Definitions

### True Positive (TP)

Correctly predicted positive

### True Negative (TN)

Correctly predicted negative

### False Positive (FP)

Incorrectly predicted positive (Type I Error)

### False Negative (FN)

Incorrectly predicted negative (Type II Error)

# Classification Metrics

---

## Accuracy

Formula:  $(TP + TN) / (TP + TN + FP + FN)$

*Use when: Balanced datasets*

⚠ Misleading for imbalanced data

## Precision

Formula:  $TP / (TP + FP)$

*Use when: Cost of FP is high*

⚠ Ignores false negatives

## Recall (Sensitivity)

Formula:  $TP / (TP + FN)$

*Use when: Cost of FN is high*

⚠ Ignores false positives

## F1-Score

Formula:  $2 \times (\text{Precision} \times \text{Recall}) / (\text{Precision} + \text{Recall})$

*Use when: Need balance between precision & recall*

⚠ Equal weight to both metrics

## Specificity

Formula:  $TN / (TN + FP)$

*Use when: True negative rate matters*

⚠ Complements sensitivity

## ROC-AUC

Formula: Area Under ROC Curve

*Use when: Overall model discrimination ability*

⚠ Not suitable for imbalanced data

# Regression Metrics

## Mean Absolute Error (MAE)

$$MAE = (1/n) \sum |y_i - \hat{y}_i|$$

Average absolute difference between predicted and actual

- ✓ Easy to interpret, same units as target ⚠ Doesn't penalize large errors heavily
- $[0, \infty)$ , Lower is better

## Mean Squared Error (MSE)

$$MSE = (1/n) \sum (y_i - \hat{y}_i)^2$$

Average of squared differences

- ✓ Penalizes large errors more ⚠ Units are squared, sensitive to outliers
- $[0, \infty)$ , Lower is better

## Root Mean Squared Error (RMSE)

$$RMSE = \sqrt{MSE}$$

Square root of MSE, same units as target

- ✓ Interpretable, penalizes large errors ⚠ Sensitive to outliers
- $[0, \infty)$ , Lower is better

## R<sup>2</sup> Score (Coefficient of Determination)

$$R^2 = 1 - (SS_{res} / SS_{tot})$$

Proportion of variance explained by model

- ✓ Normalized, easy to interpret ⚠ Can be negative for poor models
- $(-\infty, 1]$ , Higher is better, 1 = perfect

## Mean Absolute Percentage Error (MAPE)

$$MAPE = (100/n) \sum |(y_i - \hat{y}_i) / y_i|$$

Average percentage error

- ✓ Scale-independent, interpretable ⚠ Undefined when  $y_i = 0$ , biased toward low predictions
- $[0, \infty)$ , Lower is better

## Adjusted R<sup>2</sup>

$$Adj R^2 = 1 - [(1-R^2)(n-1)/(n-p-1)]$$

R<sup>2</sup> adjusted for number of predictors

- ✓ Accounts for model complexity ⚠ More complex to calculate
- $(-\infty, 1]$ , Higher is better

# Cross-Validation

## Why Cross-Validation?

- Maximizes use of limited data
- More reliable performance estimates
- Reduces variance in evaluation
- Helps detect overfitting

### K-Fold Cross-Validation

1. Split data into K equal folds
2. For each fold i:
  - Train on K-1 folds
  - Validate on fold i
3. Average performance across K folds

Typical K: 5 or 10

*Use: Standard approach for most problems*

### Stratified K-Fold

1. Maintain class distribution in each fold
2. Ensures each fold is representative
3. Same process as K-Fold
4. Especially important for imbalanced data

*Use: Classification with imbalanced classes*

### Leave-One-Out (LOO)

1. K = n (number of samples)
2. Each sample used once as validation
3. Very computationally expensive
4. Nearly unbiased estimator

*Use: Small datasets, when computational cost is acceptable*

### Time Series Split

1. Respects temporal ordering
2. Train on past, test on future
3. No shuffling of data
4. Prevents data leakage

*Use: Time-dependent data (stock prices, weather)*

# ROC Curve and AUC

---

## What is ROC Curve?

Plot of True Positive Rate (TPR) vs False Positive Rate (FPR) at different classification thresholds

**TPR (Recall/Sensitivity):**

$$\text{TPR} = \text{TP} / (\text{TP} + \text{FN})$$

**FPR:**

$$\text{FPR} = \text{FP} / (\text{FP} + \text{TN})$$

## AUC (Area Under Curve)

- AUC = 1.0: Perfect classifier
- AUC = 0.9-1.0: Excellent
- AUC = 0.8-0.9: Good
- AUC = 0.7-0.8: Fair
- AUC = 0.5-0.7: Poor
- AUC = 0.5: Random classifier
- AUC < 0.5: Worse than random

## Advantages

- Threshold-independent evaluation
- Good for comparing models
- Works well for balanced datasets
- Visualizes trade-off between TPR and FPR
- Single metric (AUC) summarizes performance

## Limitations

- ⚠ Misleading for highly imbalanced datasets
- ⚠ Doesn't reflect real-world costs of errors
- ⚠ May hide poor performance on minority class

*Alternative for Imbalanced Data: Precision-Recall Curve*

# Bias-Variance Tradeoff

## Underfitting

*High Bias*

### Characteristics:

- Model too simple
- Fails to capture patterns
- Poor training performance
- Poor test performance
- High training error
- High validation error

### Solutions:

- Increase model complexity
- Add more features
- Reduce regularization
- Train longer

## Good Fit

*Optimal Balance*

### Characteristics:

- Right model complexity
- Captures true patterns
- Good training performance
- Good test performance
- Low training error
- Low validation error

### Solutions:

- This is the goal!
- Monitor performance
- Continue validation
- Deploy with confidence

## Overfitting

*High Variance*

### Characteristics:

- Model too complex
- Memorizes training data
- Excellent training performance
- Poor test performance
- Very low training error
- High validation error

### Solutions:

- Reduce model complexity
- Add more training data
- Use regularization
- Early stopping

$$\text{Total Error} = \text{Bias}^2 + \text{Variance} + \text{Irreducible Error}$$

# Machine Learning Best Practices

---

## Data Management

- Split before preprocessing
- Use stratified splits
- Document sources
- Handle missing values

## Model Development

- Start simple
- Establish baseline
- Use cross-validation
- Track experiments

## Evaluation & Deployment

- Align metrics to goals
- Test on held-out data
- Monitor data drift
- Document assumptions

# Common Pitfalls to Avoid

---

## Data Leakage

*Example: Preprocessing before split*

## Look-Ahead Bias

*Example: Using future information*

## Selection Bias

*Example: Non-representative data*

## Class Imbalance

*Example: 99% negative, 1% positive*

## Wrong Metrics

*Example: Accuracy for imbalanced data*

## Overfitting to Validation

*Example: Excessive tuning*

# Key Takeaways

---

## Machine Learning = Learning from Data

- Algorithms discover patterns without explicit programming

## Choose the Right Algorithm

- No single algorithm is best for all problems
  - Linear/Logistic Regression for simple relationships
  - Decision Trees/Random Forests for complex, non-linear data

## Balance Bias and Variance

- Too simple → High Bias (underfitting)
- Too complex → High Variance (overfitting)

## Evaluate Properly

- Use appropriate metrics for your problem
- Always validate on unseen test data

# Thank you!



UNIVERSITY OF  
**CALGARY**